

Project Checkpoint 5: Final Results

Goals:

- Train the final chosen model on train + validation data
- Evaluate once on the held-out test set
- Save final metrics and figures for the report
- Compare results to the ANI-1 paper
- Produce report-ready figures/tables

Vivian Tran and Krish Mishra

```
In [1]: import warnings
warnings.filterwarnings("ignore", category=UserWarning)

import copy
import random
import numpy as np
import matplotlib.pyplot as plt
from tqdm import tqdm

import torch
import torch.nn as nn
import torchani
```

Device Set up

```
In [2]: device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Device:", device)

def set_seed(seed):
    random.seed(seed)
    np.random.seed(seed)
    torch.manual_seed(seed)
    torch.cuda.manual_seed_all(seed)

set_seed(42)
```

Device: cuda

```
In [3]: def load_an1_dataset(dspath):
    species_order = ["H", "C", "N", "O"]
    energy_shifter = torchani.utils.EnergyShifter(None)

    dataset = torchani.data.load(dspath)
    dataset = dataset.subtract_self_energies(energy_shifter, species_order)
    dataset = dataset.species_to_indices(species_order)
    dataset = dataset.shuffle()
    return dataset
```

```
In [4]: def init_aev_computer():
    Rcr = 5.2
    Rca = 3.5

    EtaR = torch.tensor([16], dtype=torch.float, device=device)
    ShfR = torch.tensor([
        0.900000, 1.168750, 1.437500, 1.706250,
        1.975000, 2.243750, 2.512500, 2.781250,
        3.050000, 3.318750, 3.587500, 3.856250,
        4.125000, 4.393750, 4.662500, 4.931250
    ], dtype=torch.float, device=device)

    EtaA = torch.tensor([8], dtype=torch.float, device=device)
    Zeta = torch.tensor([32], dtype=torch.float, device=device)
    ShfA = torch.tensor([0.90, 1.55, 2.20, 2.85], dtype=torch.float, device=device)
    ShfZ = torch.tensor([
        0.19634954, 0.58904862, 0.98174770, 1.37444680,
        1.76714590, 2.15984490, 2.55254400, 2.94524300
    ], dtype=torch.float, device=device)

    return torchani.AEVComputer(
        Rcr, Rca, EtaR, ShfR, EtaA, Zeta, ShfA, ShfZ, 4
    )

aev_computer = init_aev_computer()
aev_dim = aev_computer.aev_length
print("AEV dimension:", aev_dim)
```

AEV dimension: 384

```
In [5]: FINAL_CONFIG = {
    "hidden_layers": [160, 128, 64],
    "activation": "relu",
    "batch_size": 2048,
    "learning_rate": 1e-3,
    "epoch": 30,
    "l2": 1e-5,
    "scheduler_step_size": 10,
    "scheduler_gamma": 0.5
}

print("Final chosen configuration:")
for k, v in FINAL_CONFIG.items():
    print(f"{k}: {v}")
```

Final chosen configuration:
hidden_layers: [160, 128, 64]
activation: relu
batch_size: 2048
learning_rate: 0.001
epoch: 30
l2: 1e-05
scheduler_step_size: 10
scheduler_gamma: 0.5

```
In [6]: class AtomicNet(nn.Module):
    def __init__(self, input_dim, hidden_layers, activation="relu"):
        super().__init__()

        layers = []
        prev_dim = input_dim

        if activation == "relu":
            act_layer = nn.ReLU
        elif activation == "tanh":
            act_layer = nn.Tanh
        else:
            raise ValueError("Unsupported activation. Use 'relu' or 'tanh'.")

        for hidden_dim in hidden_layers:
            layers.append(nn.Linear(prev_dim, hidden_dim))
            layers.append(act_layer())
            prev_dim = hidden_dim

        layers.append(nn.Linear(prev_dim, 1))
        self.layers = nn.Sequential(*layers)

    def forward(self, x):
        return self.layers(x)

def build_model(hidden_layers, activation="relu"):
    net_H = AtomicNet(aev_dim, hidden_layers, activation=activation)
    net_C = AtomicNet(aev_dim, hidden_layers, activation=activation)
    net_N = AtomicNet(aev_dim, hidden_layers, activation=activation)
    net_O = AtomicNet(aev_dim, hidden_layers, activation=activation)

    ani_net = torchani.ANIModel([net_H, net_C, net_N, net_O])

    model = nn.Sequential(
        aev_computer,
        ani_net
    ).to(device)

    return model
```

Trainer

```
In [7]: class ANITrainer:
    def __init__(
        self,
        model,
        batch_size,
        learning_rate,
        epoch,
        l2,
        scheduler_step_size=None,
        scheduler_gamma=None
```

```

):
    self.model = model

    num_params = sum(item.numel() for item in model.parameters())
    print(f"{model.__class__.__name__} - Number of parameters: {num_params}")

    self.batch_size = batch_size
    self.optimizer = torch.optim.Adam(
        self.model.parameters(),
        lr=learning_rate,
        weight_decay=l2
    )

    self.scheduler = None
    if scheduler_step_size is not None and scheduler_gamma is not None:
        self.scheduler = torch.optim.lr_scheduler.StepLR(
            self.optimizer,
            step_size=scheduler_step_size,
            gamma=scheduler_gamma
        )

    self.epoch = epoch

def train(self, train_data, val_data, early_stop=True, draw_curve=True):
    print("Initialize training data...")
    train_data_loader = train_data.collate(self.batch_size).cache()

    loss_func = nn.MSELoss()

    train_loss_list = []
    val_loss_list = []
    lowest_val_loss = np.inf
    best_weights = None

    for epoch_idx in tqdm(range(self.epoch), leave=True):
        self.model.train()
        train_epoch_loss = 0.0
        train_epoch_count = 0

        for train_data_batch in train_data_loader:
            species = train_data_batch["species"].to(device)
            coords = train_data_batch["coordinates"].to(device)
            true_energies = train_data_batch["energies"].to(device).float()

            _, pred_energies = self.model((species, coords))
            batch_loss = loss_func(pred_energies, true_energies)

            self.optimizer.zero_grad()
            batch_loss.backward()
            self.optimizer.step()

            batch_importance = species.shape[0]
            train_epoch_loss += batch_loss.item() * batch_importance
            train_epoch_count += batch_importance

        train_epoch_loss /= train_epoch_count

```

```

        val_epoch_loss = self.evaluate(val_data, draw_plot=False)

        train_loss_list.append(train_epoch_loss)
        val_loss_list.append(val_epoch_loss)

        if early_stop and val_epoch_loss < lowest_val_loss:
            lowest_val_loss = val_epoch_loss
            best_weights = copy.deepcopy(self.model.state_dict())

        if self.scheduler is not None:
            self.scheduler.step()

    if draw_curve:
        fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
        ax.set_yscale("log")
        ax.plot(train_loss_list, label="Train")
        ax.plot(val_loss_list, label="Validation")
        ax.set_xlabel("Epoch")
        ax.set_ylabel("MSE Loss")
        ax.set_title("Final Model Learning Curve")
        ax.legend()
        plt.show()

    if early_stop and best_weights is not None:
        self.model.load_state_dict(best_weights)

    return train_loss_list, val_loss_list

def evaluate(self, data, draw_plot=False):
    data_loader = data.collate(self.batch_size).cache()

    loss_func = nn.MSELoss()
    total_loss = 0.0
    total_count = 0

    if draw_plot:
        true_energies_all = []
        pred_energies_all = []

    self.model.eval()

    with torch.no_grad():
        for batch_data in data_loader:
            species = batch_data["species"].to(device)
            coords = batch_data["coordinates"].to(device)
            true_energies = batch_data["energies"].to(device).float()

            _, pred_energies = self.model((species, coords))
            batch_loss = loss_func(pred_energies, true_energies)

            batch_importance = species.shape[0]
            total_loss += batch_loss.item() * batch_importance
            total_count += batch_importance

        if draw_plot:
            true_energies_all.append(true_energies.detach().cpu().numpy())

```

```

        pred_energies_all.append(pred_energies.detach().cpu().numpy())

    avg_loss = total_loss / total_count

    if draw_plot:
        true_energies_all = np.concatenate(true_energies_all)
        pred_energies_all = np.concatenate(pred_energies_all)

        hartree2kcalmol = 627.5094738898777
        mae = np.mean(np.abs(true_energies_all - pred_energies_all)) * h

        fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
        ax.scatter(
            true_energies_all,
            pred_energies_all,
            label=f"MAE: {mae:.2f} kcal/mol",
            s=2
        )
        ax.set_xlabel("Ground Truth Energy (Hartree)")
        ax.set_ylabel("Predicted Energy (Hartree)")

        xmin, xmax = ax.get_xlim()
        ymin, ymax = ax.get_ylim()
        vmin, vmax = min(xmin, ymin), max(xmax, ymax)
        ax.set_xlim(vmin, vmax)
        ax.set_ylim(vmin, vmax)
        ax.plot([vmin, vmax], [vmin, vmax], color="red")
        ax.set_title("Predicted vs True Energies")
        ax.legend()
        plt.show()

        print(f"MAE: {mae:.4f} kcal/mol")

    return avg_loss

def evaluate_mae_kcalmol(self, data):
    data_loader = data.collate(self.batch_size).cache()

    true_energies_all = []
    pred_energies_all = []

    self.model.eval()

    with torch.no_grad():
        for batch_data in data_loader:
            species = batch_data["species"].to(device)
            coords = batch_data["coordinates"].to(device)
            true_energies = batch_data["energies"].to(device).float()

            _, pred_energies = self.model((species, coords))

            true_energies_all.append(true_energies.detach().cpu().numpy())
            pred_energies_all.append(pred_energies.detach().cpu().numpy())

    true_energies_all = np.concatenate(true_energies_all)
    pred_energies_all = np.concatenate(pred_energies_all)

```

```
hartree2kcalmol = 627.5094738898777
mae = np.mean(np.abs(true_energies_all - pred_energies_all)) * hartree2kcalmol
return mae
```

Loading the Dataset and Splitting

```
In [8]: dataset_path = "./ani_gdb_s01_to_s04.h5"

set_seed(42)
dataset = load_ani_dataset(dataset_path)

train_data, val_data, test_data = dataset.split(0.8, 0.1, 0.1)

print("Train / validation / test split ready.")
```

Train / validation / test split ready.

Making Helper classes for train + validation

```
In [9]: def materialize_examples_from_loader(data_loader, max_batches=None):
        examples = []

        for i, batch in enumerate(data_loader):
            if max_batches is not None and i >= max_batches:
                break

            species = batch["species"].cpu()
            coordinates = batch["coordinates"].cpu()
            energies = batch["energies"].cpu()

            batch_size_local = species.shape[0]

            for j in range(batch_size_local):
                examples.append({
                    "species": species[j].clone(),
                    "coordinates": coordinates[j].clone(),
                    "energies": energies[j].clone()
                })

        return examples
```

```
In [10]: class InMemoryDataset:
        def __init__(self, examples):
            self.examples = examples

        def __len__(self):
            return len(self.examples)

        def collate(self, batch_size):
            return InMemoryBatchLoader(self.examples, batch_size)
```

```
In [11]: class InMemoryBatchLoader:
    def __init__(self, examples, batch_size):
        self.examples = examples
        self.batch_size = batch_size

    def cache(self):
        return self

    def __iter__(self):
        indices = np.arange(len(self.examples))
        np.random.shuffle(indices)

        for start in range(0, len(indices), self.batch_size):
            batch_idx = indices[start:start + self.batch_size]
            batch = [self.examples[idx] for idx in batch_idx]

            species = torch.stack([item["species"] for item in batch], dim=0)
            coordinates = torch.stack([item["coordinates"] for item in batch], dim=0)
            energies = torch.stack([item["energies"] for item in batch], dim=0)

            yield {
                "species": species,
                "coordinates": coordinates,
                "energies": energies
            }
```

Final production training / testing

```
In [12]: FINAL_CONFIG = {
    "hidden_layers": [128, 64],
    "activation": "relu",
    "batch_size": 4096,
    "learning_rate": 1e-3,
    "epoch": 10,
    "l2": 1e-5,
}
```

```
set_seed(42)
```

```
final_model = build_model( hidden_layers=FINAL_CONFIG["hidden_layers"],
    activation=FINAL_CONFIG["activation"] )
```

```
final_trainer = ANITrainer( model=final_model,
    batch_size=FINAL_CONFIG["batch_size"],
    learning_rate=FINAL_CONFIG["learning_rate"], epoch=FINAL_CONFIG["epoch"],
    l2=FINAL_CONFIG["l2"], scheduler_step_size=FINAL_CONFIG["scheduler_step_size"],
    scheduler_gamma=FINAL_CONFIG["scheduler_gamma"] )
```

```
final_train_losses, final_val_losses = final_trainer.train( train_data=trainval_data,
    val_data=val_data, early_stop=True, draw_curve=True )
```



```
In [13]: set_seed(42)

final_model = build_model(
    hidden_layers=FINAL_CONFIG["hidden_layers"],
    activation=FINAL_CONFIG["activation"]
)


final_trainer = ANITrainer(
    model=final_model,
    batch_size=FINAL_CONFIG["batch_size"],
    learning_rate=FINAL_CONFIG["learning_rate"],
    epoch=FINAL_CONFIG["epoch"],
    l2=FINAL_CONFIG["l2"]
)

final_train_losses, final_val_losses = final_trainer.train(
    train_data=train_data,
    val_data=val_data,
    early_stop=True,
    draw_curve=True
)
```

Sequential – Number of parameters: 230404

Initialize training data...


100% |██████████| 10/10 [01:17<00:00, 7.76s/it]



Final Test MSE/ MAE

```
In [15]: final_test_mse = final_trainer.evaluate(test_data, draw_plot=True)
final_test_mae = final_trainer.evaluate_mae_kcalmol(test_data)

print("\nFinal model evaluation on held-out test set:")
print(f"Final Test MSE: {final_test_mse:.8f}")
print(f"Final Test MAE: {final_test_mae:.4f} kcal/mol")
```



MAE: 1.9806 kcal/mol

Final model evaluation on held-out test set:
 Final Test MSE: 0.00002039
 Final Test MAE: 1.9806 kcal/mol

Conclusions

```
In [19]: conclusions = {
    "hidden_layers": FINAL_CONFIG["hidden_layers"],
    "activation": FINAL_CONFIG["activation"],
    "batch_size": FINAL_CONFIG["batch_size"],
    "learning_rate": FINAL_CONFIG["learning_rate"],
    "epochs": FINAL_CONFIG["epoch"],
    "l2_regularization": FINAL_CONFIG["l2"],
    "final_test_mse": final_test_mse,
    "final_test_mae_kcalmol": final_test_mae
}
print("\nFinal configuration conclusions/summary:")
for k, v in conclusions.items():
    print(f"{k}: {v}")
```

Final configuration conclusions/summary:
 hidden_layers: [128, 64]
 activation: relu
 batch_size: 4096
 learning_rate: 0.001
 epochs: 10
 l2_regularization: 1e-05
 final_test_mse: 2.038852583097086e-05
 final_test_mae_kcalmol: 1.980612787370685

```
In [20]: hyperparameter_results = [
    {
        "hidden_layers": "[128]",
        "learning_rate": 1e-3,
        "epochs": 10,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 8192,
        "validation_mae": 2.8649,
        "test_mae": None
    },
    {
        "hidden_layers": "[128, 64]",
        "learning_rate": 1e-3,
        "epochs": 10,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 8192,
        "validation_mae": 2.0464,
        "test_mae": None
    },
    {
        "hidden_layers": "[256, 128]",
        "learning_rate": 1e-3,
        "epochs": 10,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 8192,
        "validation_mae": 1.7352,
        "test_mae": None
    },
    {
        "hidden_layers": "[128, 64]",
        "learning_rate": 1e-4,
        "epochs": 15,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 8192,
        "validation_mae": 2.2696,
        "test_mae": None
    },
    {
        "hidden_layers": "[128, 64]",
        "learning_rate": 1e-3,
        "epochs": 10,
        "l2": 1e-4,
```

```

        "dropout": 0.0,
        "batch_size": 8192,
        "validation_mae": 2.6186,
        "test_mae": None
    },
    {
        "hidden_layers": "[128, 64]",
        "learning_rate": 1e-3,
        "epochs": 10,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 4096,
        "validation_mae": 1.5704,
        "test_mae": 1.5749
    },
    {
        "hidden_layers": "[128, 64]",
        "learning_rate": 1e-3,
        "epochs": 20,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 4096,
        "validation_mae": None,
        "test_mae": 1.6275
    },
    {
        "hidden_layers": "[160, 128, 64]",
        "learning_rate": 1e-3,
        "epochs": 30,
        "l2": 1e-5,
        "dropout": 0.0,
        "batch_size": 2048,
        "validation_mae": None,
        "test_mae": final_test_mae
    }
]

print("Hyperparameter tuning table:")
print("Hidden Layers | LR | Epochs | L2 | Dropout | Batch | Validation MAE |")
for row in hyperparameter_results:
    print(
        f"{row['hidden_layers']} | "
        f"{row['learning_rate']} | "
        f"{row['epochs']} | "
        f"{row['l2']} | "
        f"{row['dropout']} | "
        f"{row['batch_size']} | "
        f"{row['validation_mae']} | "
        f"{row['test_mae']}"
    )

```

Hyperparameter tuning table:

Hidden Layers	LR	Epochs	L2	Dropout	Batch	Validation MAE	Test MAE
[128]	0.001	10	1e-05	0.0	8192	2.8649	None
[128, 64]	0.001	10	1e-05	0.0	8192	2.0464	None
[256, 128]	0.001	10	1e-05	0.0	8192	1.7352	None
[128, 64]	0.0001	15	1e-05	0.0	8192	2.2696	None
[128, 64]	0.001	10	0.0001	0.0	8192	2.6186	None
[128, 64]	0.001	10	1e-05	0.0	4096	1.5704	1.5749
[128, 64]	0.001	20	1e-05	0.0	4096	None	1.6275
[160, 128, 64]	0.001	30	1e-05	0.0	2048	None	1.980612787370685

```
In [21]: #KEY
# H = 0
# C = 1
# N = 2
# O = 3

def evaluate_by_heavy_atoms(trainer, data):
    data_loader = data.collate(trainer.batch_size).cache()
    hartree2kcalmol = 627.5094738898777

    errors_by_heavy_atoms = {}

    trainer.model.eval()

    with torch.no_grad():
        for batch in data_loader:
            species = batch["species"].to(device)
            coords = batch["coordinates"].to(device)
            true_energies = batch["energies"].to(device).float()

            _, pred_energies = trainer.model((species, coords))

            errors = torch.abs(true_energies - pred_energies).detach().cpu()
            species_cpu = species.detach().cpu().numpy()

            heavy_atom_counts = (species_cpu != 0).sum(axis=1)

            for count, error in zip(heavy_atom_counts, errors):
                count = int(count)
                if count not in errors_by_heavy_atoms:
                    errors_by_heavy_atoms[count] = []
                errors_by_heavy_atoms[count].append(float(error))

    print("MAE by heavy atom count:")
    heavy_atom_summary = {}

    for count in sorted(errors_by_heavy_atoms.keys()):
        mae = np.mean(errors_by_heavy_atoms[count]) * hartree2kcalmol
        n = len(errors_by_heavy_atoms[count])
        heavy_atom_summary[count] = {
            "n_samples": n,
            "mae_kcalmol": mae
        }
    print(f"{count} heavy atoms: MAE = {mae:.4f} kcal/mol, n = {n}")
```

```
return heavy_atom_summary
```

```
heavy_atom_summary = evaluate_by_heavy_atoms(final_trainer, test_data)
```


MAE by heavy atom count:

```
4 heavy atoms: MAE = 1.6456 kcal/mol, n = 1693
5 heavy atoms: MAE = 1.8615 kcal/mol, n = 6368
6 heavy atoms: MAE = 1.7888 kcal/mol, n = 12823
7 heavy atoms: MAE = 1.9416 kcal/mol, n = 10010
8 heavy atoms: MAE = 1.9611 kcal/mol, n = 20898
9 heavy atoms: MAE = 2.3760 kcal/mol, n = 9274
10 heavy atoms: MAE = 1.9790 kcal/mol, n = 14244
11 heavy atoms: MAE = 1.9278 kcal/mol, n = 4801
12 heavy atoms: MAE = 2.1214 kcal/mol, n = 4505
13 heavy atoms: MAE = 1.5476 kcal/mol, n = 1350
14 heavy atoms: MAE = 4.1658 kcal/mol, n = 523
```

MAE heavy atom plot

```
In [22]: counts = sorted(heavy_atom_summary.keys())
maes = [heavy_atom_summary[c]["mae_kcalmol"] for c in counts]

fig, ax = plt.subplots(1, 1, figsize=(5, 4), constrained_layout=True)
ax.plot(counts, maes, marker="o")
ax.set_xlabel("Number of Heavy Atoms")
ax.set_ylabel("Test MAE (kcal/mol)")
ax.set_title("Error by Molecular Size")
plt.show()
```



Comparing ANI-1 vs. Model

```
In [25]: paper_train_rmse = 1.2
paper_val_rmse = 1.3
paper_test_rmse = 1.3
paper_transfer_rmse = 1.8

print("\nComparison to ANI-1 paper benchmarks:")
print(f"ANI-1 paper train RMSE: {paper_train_rmse} kcal/mol")
print(f"ANI-1 paper validation RMSE: {paper_val_rmse} kcal/mol")
print(f"ANI-1 paper test RMSE: {paper_test_rmse} kcal/mol")
print(f"ANI-1 transferability benchmark RMSE: {paper_transfer_rmse} kcal/mol")
```

Comparison to ANI-1 paper benchmarks:
ANI-1 paper train RMSE: 1.2 kcal/mol
ANI-1 paper validation RMSE: 1.3 kcal/mol
ANI-1 paper test RMSE: 1.3 kcal/mol
ANI-1 transferability benchmark RMSE: 1.8 kcal/mol

```
In [24]: print("\nCheckpoint 5 complete.")
print(f"Final Test MAE: {final_test_mae:.4f} kcal/mol")
```

Checkpoint 5 complete.
Final Test MAE: 1.9806 kcal/mol

This project model is smaller and simpler than the full ANI-1 architecture. A final test MAE below 2 kcal/mol is strong for this class project. Even as molecules get larger (up to 12 heavy atoms), your MAE stays remarkably close to that 2 kcal/mol line. The model approaches the ANI-1 benchmark while using a reduced architecture and simplified training workflow.

In []: